Zuu Crew
Machine Learning
Acadamy

MINI PROJECT 02

Real Estate Intelligence Platform for
Prime Lands

Course

AI Engineer Essentials

Module

Context Engineering

Weight

15% of Total Grade

Total Points

100 (+5 Bonus)

Assessment Overview

Four parts plus report equals a production-ready RAG system evaluation.

Part 1

Web Crawling

15 pts

Tests async programming, data extraction, and production crawling practices.

Part 2

Chunking Lab

25 pts

Evaluates understanding of text splitting trade-offs and vector indexing.

Part 3

Intelligence Layers

25 pts

Assesses advanced RAG implementation (CAG + CRAG) and clean architecture.

Part 4

Performance Arena

20 pts

Measures evaluation rigor, metrics implementation, and analytical thinking.

Report

Engineering Doc

15 pts

Tests technical writing, critical analysis, and production recommendations.

Time Expected	16-21 Hours (2-3 Weeks)	Success Indicator	Production-Ready RAG Skills

Part 1: The Property Crawler

(https://www.primelands.lk)

15 Points

Playwright Implementation — 5 pts

Async crawler with proper browser lifecycle management, JavaScript rendering support, and BFS traversal.

✅ Full (5) ⚠️ Partial (3) ❌ None (0)

Data Extraction Completeness — 5 pts

All metadata fields extracted correctly (property_id, title, address, price, bedrooms, bathrooms, sqft, amenities, agent).

✅ Full (5) ⚠️ Partial (3) ❌ None (0)

Output Format Correctness — 5 pts

Both Markdown files and JSONL corpus generated with correct structure and schema.

✅ Full (5) ⚠️ Partial (3) ❌ None (0)

Red Flags

- Hardcoded URLs or paths
- No rate limiting
- Synchronous blocking calls
- Missing error handling

Quick Check

Run 01_crawl_primelands.ipynb

Verify data/primelands_corpus.jsonl has 5+ entries

Check metadata completeness

Part 2: The Chunking Lab

25 Points

All 5 Strategies Implemented — 10 pts

Semantic, Fixed, Sliding, Parent-Child, and Late Chunking functions work correctly with proper token counting.

✅ Full (10) ⚠️ Partial (6-8) ❌ Low (2-4)

Qdrant Indexing — 5 pts

Persistent index created using Qdrant, all 5 collections populated with embeddings and rich metadata.

✅ Full (5) ⚠️ Partial (3) ❌ None (0)

Comparison Metrics — 10 pts

Complete comparison table with chunk count, avg size, index size, and retrieval time for all strategies.

✅ Full (10) ⚠️ Partial (6) ❌ Low (2)

Red Flags

- Hardcoded chunk sizes (ignore config)
- No token counting (char split only)
- Missing parent-child linking
- Identical results across strategies

Quick Check

- Verify 5 JSONL files exist
- Check Qdrant folder has 5 collections
- Review comparison table for meaningful differences

Part 3: Intelligence Layers

25 Points

RAGService with LCEL — 8 pts

Modern LCEL chain (Runnable), proper retriever integration, inline citations with evidence URLs.

✅ Full (8) ⚠️ Partial (5) ❌ None (0)

CAGService Caching — 8 pts

Two-tier cache (FAQs + History), semantic similarity (cosine > 0.90), cache hit tracking, FAQ pre-warming.

✅ Full (8) ⚠️ Partial (5) ❌ None (0)

CRAGService Correction — 9 pts

Confidence scoring implemented, corrective retrieval triggers below threshold (0.6), demonstrates improvement.

✅ Full (9) ⚠️ Partial (5-6) ❌ None (0)

Red Flags

All code in notebooks (no service layer)

No architecture separation

Hardcoded prompts

Missing cache statistics

Quick Check

Test CAG with FAQ query (< 500ms)

Test CRAG with complex query (check confidence scores)

Part 4: Performance Arena

20 Points

Chunking Strategy Comparison — 8 pts

Evaluates 10 queries across 5 strategies. Must include Precision@5, Recall@5, Answer Relevance, and Latency metrics. Identifies clear winner.

CAG Cache Effectiveness — 6 pts

Simulates 100 queries. Calculates cache hit rate, latency improvement, and estimated cost savings from avoided API calls.

Red Flags

- Cherry-picked queries
- No statistical significance
- Subjective evaluation only
- Missing output files

CRAG Correction Impact — 6 pts

Compares RAG vs CRAG on 20 queries. Tracks correction frequency, confidence improvement, and answer quality gains.

BONUS: Cost Analysis — +5 pts

Detailed monthly cost breakdown (API + storage) and ROI analysis. Must use realistic assumptions for scale (e.g., 500 daily users).

Quick Check

Verify 3 output files exist:

- chunking_comparison.csv
- cag_stats.json
- crag_impact.csv

Engineering Report Grading

15 Points

Executive Summary — 3 pts

Clear winner identified, key metrics cited, concise recommendation (150-200 words).

✅ Full (3) ⚠️ Partial (2) ❌ None (0)

Methodology — 4 pts

Detailed explanation of crawling, chunking, RAG architecture, and caching strategy with justification.

✅ Full (4) ⚠️ Partial (2-3) ❌ None (0)

Analysis - The Chunking Showdown — 5 pts

Quantitative comparison table, qualitative examples, trade-off discussion, clear recommendation.

✅ Full (5) ⚠️ Partial (3) ❌ Low (1)

Conclusion & Recommendations — 3 pts

Optimal architecture identified, scalability considerations, cost projection, next steps.

✅ Full (3) ⚠️ Partial (2) ❌ None (0)

Red Flags

- Copy-pasted assignment text
- No original analysis
- Missing word count
- No citations or poor formatting

Quick Check

- PDF report 1500-2000 words
- 4 required sections present
- Includes data tables
- Demonstrates critical thinking

Code Quality Deductions

Production-ready code requires documentation, configuration management, and error handling.

Missing Docstrings — -2 pts

Functions and classes lack docstrings explaining purpose, parameters, and return values.

Fix: Add Google-style or NumPy-style docstrings

Hardcoded Values — -2 pts

Chunk sizes, API keys, file paths, or thresholds hardcoded in notebooks or scripts.

Fix: Move all configuration to config.yaml

No Error Handling — -3 pts

No try-except blocks for API calls, file operations, or network requests. Crashes on edge cases.

Fix: Add graceful error handling with messages

Other Penalties

Incomplete Submission	-5 pts
Late Submission	-10% / day
Plagiarism	0 pts

Positive Indicators

- Clean Architecture
- Comprehensive README
- Unit Tests (Bonus)
- Deployment Scripts (Bonus)